

ITCS45 Artificial Intelligence



Instructor

Teacher Thanapon Noraset

Solving OpenAI's Gym Classic environments using Hill-climbing search (sideway + restart)

By

Mr. Komsan Kongwongsupak 6288024

Miss Anyamanee Amatyakul 6288060

Mr. Kasidis Chokphaiboon 6288071

Miss Cholravee Kittimethee 6288079

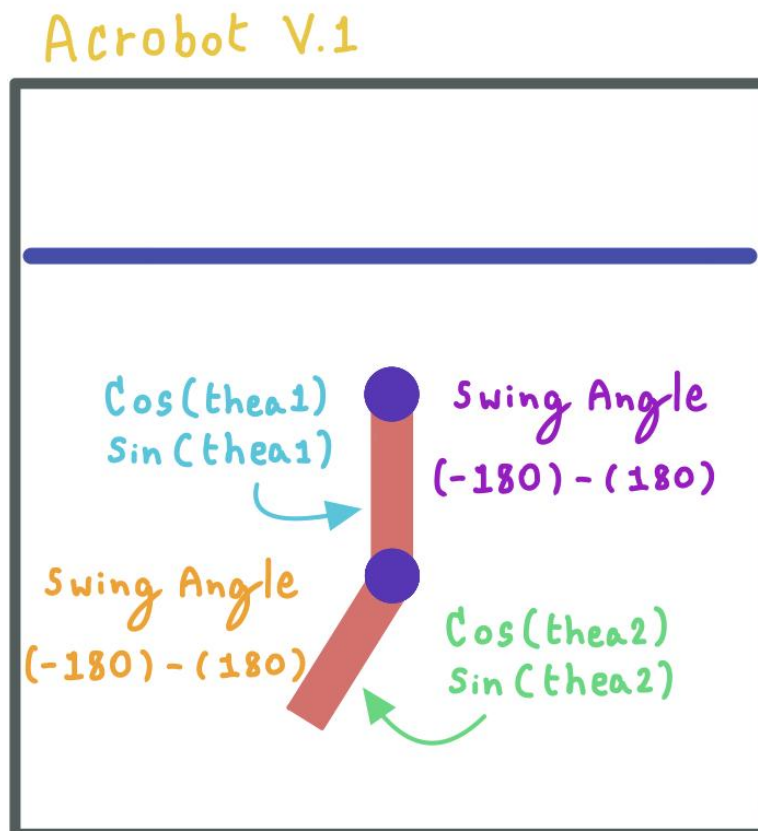
Faculty of Information and Communication Technology

Mahidol University

9 October 2021

Acrobot V1


The acrobot environment includes two joints and two links, with the joint connecting the two links being operated. Initially, the links are hung downwards, and the goal is to swing the lower links end up to a specific height. The state is made up of the $\sin()$ and $\cos()$ of the two rotating joint angles, as well as the joint angular velocities: $[\cos(\theta_1), \sin(\theta_1), \cos(\theta_2), \sin(\theta_2), V_1, V_2]$. An angle of 0 corresponds to the first link pointing downwards. The second link's angle is relative to the first link's angle. An angle of 0 denotes that the two links have the same angle. A condition of $[1, 0, 1, 0, \dots, \dots]$ indicates that both connections are pointing downward. If Sine is negative, the angle will be negative and shift left, but it will shift right if the angle is positive. Plus, if it is true, it will receive 0, and if it is false, it will receive -1 as a reward.



CartPole V0

A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track. Controlling the mechanism is applying a force of +1 or -1 to the cart. The goal is to keep the pendulum upright and from falling over. Every timestep that the pole remains upright results in a +1 reward. The episode terminates when the pole is more than 12 degrees from vertical or the cart is more than 2.4 units from the center. The state is made up of the car's position, the car's velocity, the angle of the pole, and velocity at the tip of the pole.


Cart Pole - V1



MountainCar V0

A car is on a one-dimensional track, positioned between two mountains. The goal is to drive up the mountain on the right; however, the car has no power to drive up the mountain at the one time. Therefore, it needs to drive back and forth to build up momentum. The less time it spends to reach the peak, the greater reward it will receive. The state is made up of the position of the car and the car's velocity.


Mountain Car - v0



Pendulum V0

The inverted pendulum swing-up problem is a classic problem in the control literature. The pendulum starts in a random position, and the goal is to swing it up so that it stays upright. The state is made up of $[\cos(\theta), \sin(\theta), \text{velocity}]$.

Pendulum - v0



Hill-Climbing Search Algorithm

```
145 def hillclimb_sideway(env, agent, max_iters=10000, sideways_limit=10):
146     """
147     Run a hill-climbing search, and return the final agent.
148
149     Parameters
150     -----
151     env : OpenAI Gym Environment.
152         A cart-pole environment for the agent.
153     agent : CPAgent
154         An initial agent.
155     max_iters: int
156         Maximum number of iterations to search.
157     sideways_limit
158         Number of sideways move to make before terminating.
159         Note that the sideways count reset after a new better neighbor
160         has been found.
161
162     Returns
163     -----
164     final_agent : CPAgent
165         The final agent.
166     history : List[float]
167         A list containing the scores of the best neighbors of
168         all iterations. It must include the last one that causes
169         the algorithm to stop.
170
171     """
172     cur_agent = agent
173     cur_r = simulate(env, [agent])[0]
174
175     explored = set()
176     explored.add(cur_agent)
177     history = [cur_r]
178
179     for __ in range(max_iters):
180         env.render()
181         # TODO 1: Implement hill climbing search with sideways move.
182         neighbors = cur_agent.neighbors()
183         _n = []
184         for a in neighbors:
185             if a not in explored: # we do not want to move to previously explored ones.
186                 _n.append(a)
187         neighbors = _n
188         rewards = simulate(env, neighbors)
189         best_i = np.argmax(rewards)
190         history.append(rewards[best_i])
191         if rewards[best_i] < cur_r:
192             return cur_agent, history
193         # Sideway move
194         elif rewards[best_i] == cur_r:
195             for __ in range(sideways_limit):
196                 rewards = simulate(env, neighbors)
197                 equal_i = np.argmax(rewards)
198                 history.append(rewards[equal_i])
199                 if rewards[equal_i] < cur_r:
200                     best_i = equal_i
201                     break
202             return cur_agent, history
203         #
204         cur_agent = neighbors[best_i]
205         cur_r = rewards[best_i]
206
207     # pass
208     return cur_agent, history
209
```

According to the code in Hillclimb, it finds the highest neighbor and returns that agent. However, in the hillclimb_sideway, if it finds a higher neighbor, it will stop and walk sideways a few times. If the program walks sideways ten-time but cannot find the higher neighbor, it will return to the current agent.

```
210
211 def hillclimb_restart(env, agent):
212     """Run a hill-climbing search, and return the final agent."""
213     best_agent, rewards = hillclimb_sideway(env, agent)
214     best_reward = max(rewards)
215     for __ in range(30):
216         cur_agent = ACAgent()
217         temp_agent, history = hillclimb_sideway(env, cur_agent)
218         reward = max(history)
219         if best_reward < reward:
220             best_agent = temp_agent
221             best_reward = reward
222
223     return best_agent, history
224
```

In hillclimbing_restart, we run time first to collect the initial reward of the initial agent into our rewards log. Then we will run thirty more times, find the best result, and collect the higher agent into our log until we find the best agent. (It will not be the same agent) P.S. We run code thirty times to guarantee the goal.

Testing result

Acrobot (input-output-average)

	Cos (theta1)	Sin (theta1)	Cos (theta2)	Sin (theta2)	Velocity 1	Velocity 2	Bias	Rewards
Round 1	-0.0289	0.072	0.0391	-0.115	0.076	-0.000841	-0.0208	-298
Round 2	-0.000721	0.0536	0.00617	0.119	-0.083	-0.0434	-0.03	-298
Round 3	0.0332	0.0674	-0.163	-0.072	0.055	0.00925	-0.152	-298
Round 4	-0.217	-0.019	0.00338	-0.0812	-0.0149	0.0485	-0.049	-298
Round 5	-0.154	0.0454	0.0349	0.0571	0.0142	0.0286	-0.0451	-298
Average								-298

Cartpole (input-output-average)

	Cart position	Cart velocity	Pole angle	Pole velocity at tip	Bias	Rewards
Round 1	0.0999	-0.0877	0.085	0.0586	0.00779	149
Round 2	0.0565	0.0608	0.0865	0.686	-0.000163	500
Round 3	0.0113	-0.0268	0.134	0.105	0.00519	500
Round 4	0.0268	0.00757	0.237	0.152	-0.00223	500
Round 5	0.005	0.0757	0.154	0.00582	-0.0174	500
Average						429.8

MountainCar (input-output-average)

	Car position	Car velocity	Bias	Rewards
Round 1	0.209	0.127	0.0949	-120
Round 2	0.0354	-0.0331	-0.137	-120
Round 3	0.112	0.141	0.0778	-120
Round 4	0.0562	0.0141	-0.122	-120
Round 5	0.236	0.0377	-0.141	-120
Average				-120

Pendulum (input-output-average)

	$\cos(\theta)$	$\sin(\theta)$	Velocity	Bias	Rewards
Round 1	0.232	-0.0906	0.0237	-0.11	-568.33
Round 2	-0.17	-0.141	0.00426	-0.182	-617.69
Round 3	0.0631	-0.146	0.0258	-0.107	-566.45
Round 4	0.0473	0.0811	0.0266	-0.13	-500.71
Round 5	-0.0667	-0.113	0.0143	-0.0613	-623.87
Average					-575.41

Credits

The original code: Aj.Thanapon Noraset

Implemented code: Mr.Komsan Kongwongsupak and Miss Anyamanee Amatyakul

Environments: OpenAI Gym ([Gym \(openai.com\)](https://openai.com/gym))

Hill-Climbing Search Sideway and Random-Restart Algorithm: Mr.Komsan Kongwongsupak and Miss Anyamanee Amatyakul

Description and Testing: Mr.Kasidis Chokphaiboon and Miss Cholravee Kittimethee